

# **Parallel and Distributed Storage**

# Introduction

- Parallel machines have become quite common and affordable
  - prices of microprocessors, memory and disks have dropped sharply
- Data storage needs are growing increasingly large
  - user data at web-scale
    - 100's of millions of users, petabytes of data
  - transaction data are collected and stored for analysis.
  - multimedia objects like images/videos
- Parallel storage system requirements
  - storing large volumes of data
  - processing time-consuming decision-support queries
  - providing high throughput for transaction processing
  - Very high demands on **scalability** and **availability**

# Parallel/Distributed Data Storage History

- 1980/1990s
  - Distributed database systems with tens of nodes
- 2000s:
  - Distributed file systems with 1000s of nodes
    - Millions of Large objects (100's of megabytes)
    - Web logs, images, videos, ...
    - Typically create/append only
  - Distributed data storage systems with 1000s of nodes
    - Billions to trillions of smaller (kilobyte to megabyte) objects
    - Social media posts, email, online purchases, ...
    - Inserts, updates, deletes
  - **Key-value stores**
- 2010s: Distributed database systems with 1000s of nodes

# I/O Parallelism

- Reduce the time required to retrieve relations from disk by partitioning the relations on *multiple disks*, on *multiple nodes* (computers)
  - Our description focuses on parallelism across nodes
  - Same techniques can be used across disks on a node
- **Horizontal partitioning** – tuples of a relation are divided among many nodes such that some subset of tuple resides on each node.
  - Contrast with **vertical partitioning**, e.g.  $r(A,B,C,D)$  with primary key  $A$  into  $r1(A,B)$  and  $r2(A,C,D)$
  - By default, the word partitioning refers to horizontal partitioning

# I/O Parallelism

- Partitioning techniques (number of nodes =  $n$ ):

## Round-robin:

Send the  $i^{\text{th}}$  tuple inserted in the relation to node  $i \bmod n$ .

In a data center,


Nodes N0...N19 in Rack 1

Nodes N20...N39 in Rack 2

Total nodes  $N = 40$

NID-Person relation has tuples  
T0, T1, T2, T3 ..... T39999

**Question 4-1:** Explain how tuples will be distributed into the nodes using round-robin technique.



(e) tree-like topology

- Partitioning techniques (number of nodes =  $n$ ):

### Hash partitioning:

- Choose one or more attributes as the partitioning attributes.
- Choose hash function  $h$  with range  $0 \dots n - 1$
- Let  $i$  denote result of hash function  $h$  applied to the partitioning attribute value of a tuple. Send tuple to node  $i$ .

In a data center,

Nodes  $N_0 \dots N_{19}$  in Rack 1

Nodes  $N_{20} \dots N_{39}$  in Rack 2

Total nodes  $N = 40$


NID-Person(NID, name, DOB, street, city, district)

The tuples are:

$T_0, T_1, T_2, T_3 \dots T_{39999}$

**Question 4-2:** Explain how tuples will be distributed into the nodes using hash partitioning technique using

- NID
- Street, city, district.



(e) tree-like topology

# Range Partitioning

## Data Partitioning among Nodes

### Range partitioning technique:

- Choose an attribute as the partitioning attribute.
- A partitioning vector  $[v_0, v_1, \dots, v_{n-2}]$  is chosen.
- Let  $v$  be the partitioning attribute value of a tuple.
- Tuples such that  $v_i \leq v < v_{i+1}$  go to node  $i+1$ .
- Tuples with  $v < v_0$  go to node 0 and tuples with  $v \geq v_{n-2}$  go to node  $n-1$ .
  - E.g., with a partitioning vector  $[5, 11]$
  - a tuple with partitioning attribute value of 2 will go to node 0,
  - a tuple with value 8 will go to node 1, while
  - a tuple with value 20 will go to node 2.

In a data center,

Nodes N0...N9 in Rack 1

Nodes N20...N19 in Rack 2

Total nodes  $N = 20$

Student(Id, name, DOB, street, city, district)

The tuples are:

T0, T1, T2, T3 ..... T399

Id ranges from 202105001 to 202105400

### Question 4-3

- a. Design a partition vector on Id for storage of student relation
- b. Find partitions P0, P1, P19

# Range Partitioning

## Data Partitioning among Nodes

**Example:** There are 4 nodes: N1, N2, N3 and N4 in a parallel database system. The schema of the relation person is as follows:

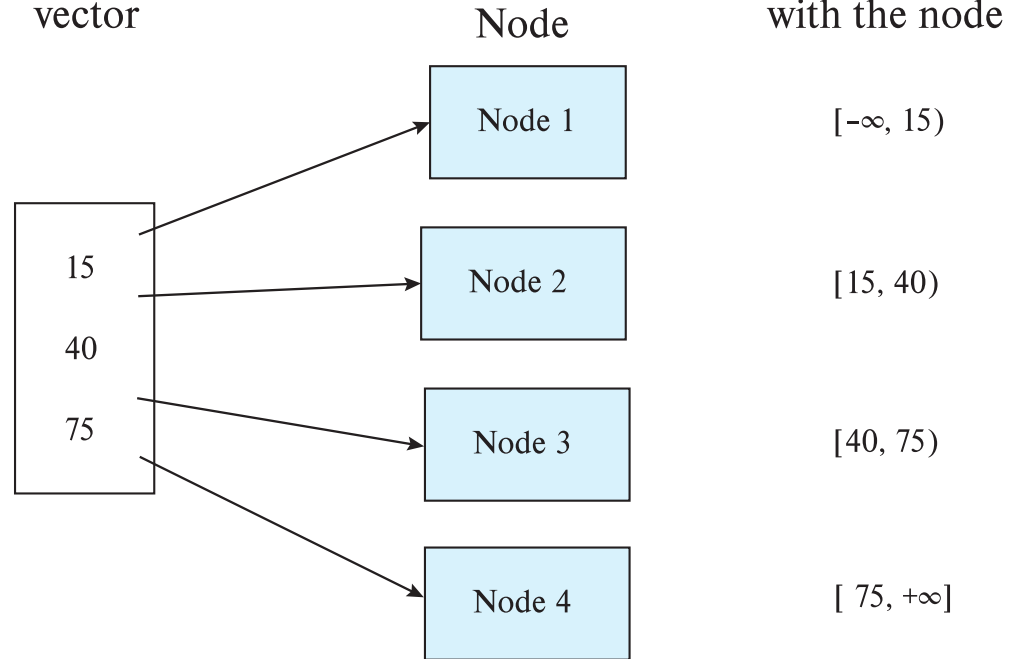
Person(NID, Name, Thana, District, Age)

Partition vector

$P[f\text{-age}] = [30, 50, 75]$

Perform range partitioning of person relation into 4 nodes.

Range partitioning  
vector



### Partitions:

Partition 1, age  $< 30$  in Node N1

Partition 2,  $30 \leq \text{age} < 50$  in Node N2

Partition 3,  $50 \leq \text{age} < 75$  in Node N3

Partition 4, age  $\geq 75$  in Node N4



# Range Partitioning

## Data Partitioning among Nodes

**Example:** There are 4 nodes: N1, N2, N3 and N4 in a parallel database system. The schema of the relation **parents** is as follows:


parents(m-NID, f-NID, c-NID, f-age)

Partition vector

$P[f\text{-age}] = [30, 50, 75]$

Perform range partitioning of parents relation into 4 nodes.

Range partitioning  
vector



Range associated  
with the node

$[-\infty, 15)$

$[15, 40)$

$[40, 75)$

$[75, +\infty]$

### Partitions:

Partition 1, f-age < 30 in Node N1

Partition 2,  $30 \leq f\text{-age} < 50$  in Node N2

Partition 3,  $50 \leq f\text{-age} < 75$  in Node N3

Partition 4, f-age  $\geq 75$  in Node N4

## Data Partitioning among Nodes

### Partitions:

#### Person relation

Partition 1, age <30 in Node N1

Partition 2, 30 ≤ age <50 in Node N2

Partition 3, 50 ≤ age <75 in Node N3

Partition 4, age ≥75 in Node N4

### Partitions:

#### Parents relation


Partition 1, f-age <30 in Node N1

Partition 2, 30 ≤ f-age <50 in Node N2

Partition 3, 50 ≤ f-age <75 in Node N3

Partition 4, f-age ≥75 in Node N4

Range partitioning  
vector



Range associated  
with the node

$[-\infty, 15)$

$[15, 40)$

$[40, 75)$

$[75, +\infty]$

Person(NID, Name, Thana, District, Age), parents(m-NID, f-NID, c-NID, f-age)  
process the query1: SELECT \* FROM person r, parents s WHERE r.age = s.f-age

#### Question 4-4:

- Explain how query 1 will be executed using 4 nodes?
- Comments on speed-up and scale-up for the system.

# Comparison of Partitioning Techniques

- Evaluate how well partitioning techniques support the following types of data access:

1. Scanning the entire relation.

2. Locating a tuple associatively – **point queries**.

- E.g.,  $r.A = 25$ .

3. Locating all tuples such that the value of a given attribute lies within a specified range – **range queries**.

- E.g.,  $10 \leq r.A < 25$ .

Select \* from person

Select \* from person where NID = 1234567890

Select \* from person where NID > 1234567890

**Question5-1:** Explain the performance of Round robin partitioning technique for

- Scanning the entire relation
- Point query
- Range query

# Comparison of Partitioning Techniques (Cont.)

## Round robin:

- Best suited for sequential scan of entire relation on each query.
  - All nodes have almost an equal number of tuples; retrieval work is thus well balanced between nodes.
- All queries must be processed at all nodes

# Comparison of Partitioning Techniques

- Evaluate how well partitioning techniques support the following types of data access:

1. Scanning the entire relation.

2. Locating a tuple associatively – **point queries**.

- E.g.,  $r.A = 25$ .

3. Locating all tuples such that the value of a given attribute lies within a specified range – **range queries**.

- E.g.,  $10 \leq r.A < 25$ .

Select \* from person

Select \* from person where NID = 1234567890

Select \* from person where NID > 1234567890

**Question5-2:** Explain the performance of hash partitioning technique for

- Scanning the entire relation
- Point query
- Range query

# Comparison of Partitioning Techniques (Cont.)

## Hash partitioning:

- Good for sequential access
  - Assuming hash function is good, and partitioning attributes form a key, tuples will be equally distributed between nodes
- Good for point queries on partitioning attribute
  - Can lookup single node, leaving others available for answering other queries.
- Range queries inefficient, must be processed at all nodes

# Comparison of Partitioning Techniques

- Evaluate how well partitioning techniques support the following types of data access:

1. Scanning the entire relation.

2. Locating a tuple associatively – **point queries**.

- E.g.,  $r.A = 25$ .

3. Locating all tuples such that the value of a given attribute lies within a specified range – **range queries**.

- E.g.,  $10 \leq r.A < 25$ .

Select \* from person

Select \* from person where NID = 1234567890

Select \* from person where NID > 1234567890

**Question5-3:** Explain the performance of range partitioning technique for

- Scanning the entire relation
- Point query
- Range query

# Comparison of Partitioning Techniques (Cont.)

## Range partitioning:

- Provides data clustering by partitioning attribute value.
  - Good for sequential access
  - Good for point queries on partitioning attribute: only one node needs to be accessed.
- For range queries on partitioning attribute, one to a few nodes may need to be accessed
  - Remaining nodes are available for other queries.
  - Good if result tuples are from one to a few blocks.
  - But if many blocks are to be fetched, they are still fetched from one to a few nodes, and potential parallelism in disk access is wasted
    - Example of **execution skew**.



# Handling Small Relations

- Partitioning not useful for small relations which fit into a single disk block or a small number of disk blocks
  - Instead, assign the relation to a single node, or
  - Replicate relation at all nodes
- For medium sized relations, choose how many nodes to partition across based on size of relation
- Large relations typically partitioned across all available nodes.

# Types of Skew

- **Data-distribution skew:** some nodes have many tuples, while others may have fewer tuples. Could occur due to
  - **Attribute-value skew.**
    - Some partitioning-attribute values appear in many tuples
      - E.g., partitioning on age in ecommerce database  
Game related DB,  $P[\text{age}] = [10, 20, 30, 40, 50, 60, 70]$   
Comments on the partition vector???
    - All the tuples with the same value for the partitioning attribute end up in the same partition.
    - Can occur with range-partitioning and hash-partitioning.

# Types of Skew

- **Partition skew.**
  - Imbalance, even without attribute –value skew
  - Badly chosen range-partition vector may assign too many tuples to some partitions and too few to others.
  - Less likely with hash-partitioning

# Types of Skew (Cont.)

- Note that **execution skew** can occur even without data distribution skew
  - E.g. relation range-partitioned on date, and most queries access tuples with recent dates
- Data-distribution skew can be avoided with range-partitioning by creating **balanced range-partitioning vectors**
- We assume for now that partitioning is **static**, that is partitioning vector is created once and not changed
  - Any change requires **repartitioning**
  - **Dynamic partitioning** once allows partition vector to be changed in a continuous manner
    - More on this later

**Question 6-1:** Explain the various reasons and types of skews in

- a. Hash partitioning
- b. Range partitioning

# Handling Skew in Range-Partitioning

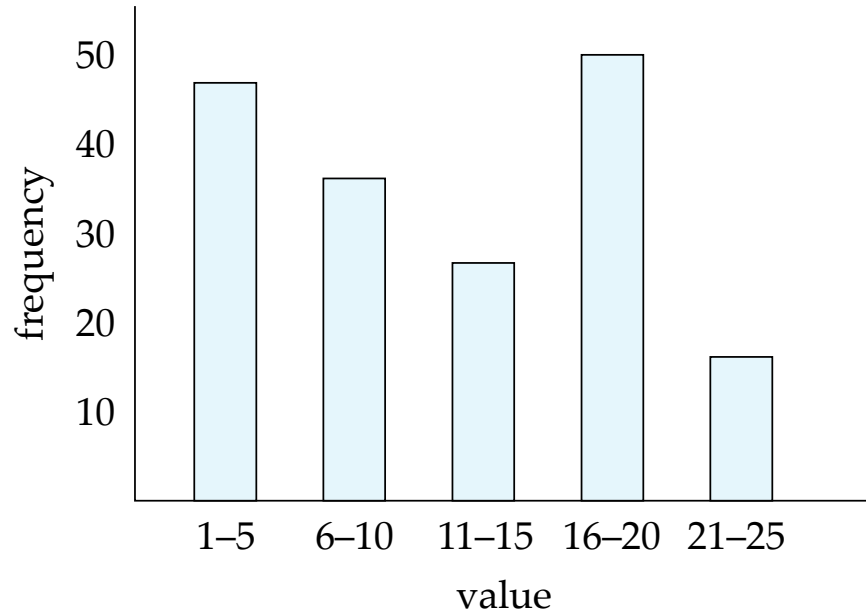
- To create a **balanced partitioning vector**
  - Sort the relation on the partitioning attribute.
  - Construct the partition vector by scanning the relation in sorted order as follows.
    - After every  $1/n^{th}$  of the relation has been read, the value of the partitioning attribute of the next tuple is added to the partition vector.
  - $n$  denotes the number of partitions to be constructed.
  - Imbalances can result if duplicates are present in partitioning attributes.

# Handling Skew in Range-Partitioning

- To reduce cost
  - Partitioning vector can be created using a random sample of tuples
  - Alternatively histograms can be used to create the partitioning vector

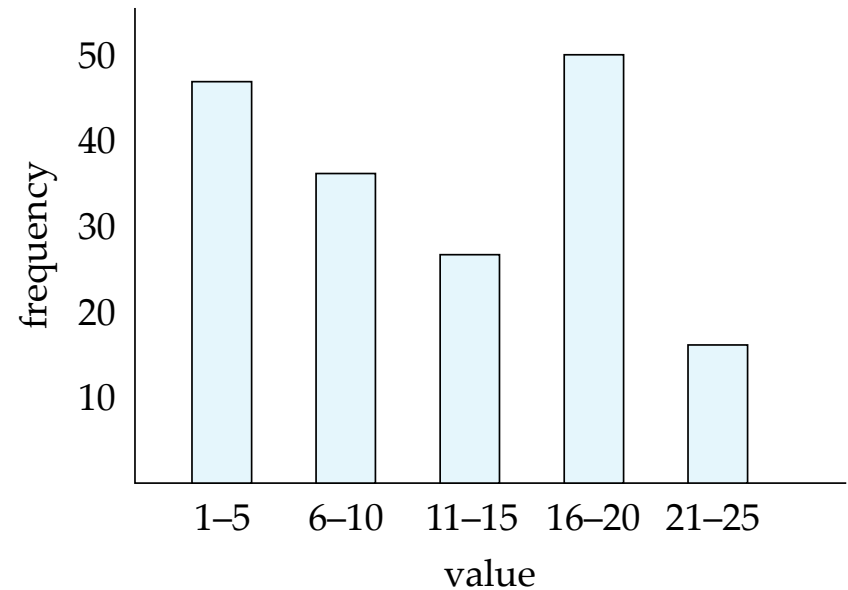
# Histograms

- Histogram on attribute *age* of relation *person*



- **Equi-width** histograms
- **Equi-depth** histograms
  - break up range such that each range has (approximately) the same number of tuples
  - E.g. (4, 8, 14, 19)
- Assume uniform distribution within each range of the histogram
- Create partitioning vector for required number of partitions based on histogram

# Histograms



## Question 6-2:

- What is the type of the given histogram?
- Define an approximate partition vector using the histogram for a parallel system with nodes N0, N1, N2 and N3.
- Draw the approximate histogram of your partition vector.