

# **Concurrency Control in Centralized Databases**

# Atomocity Using Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
  - *exclusive* (X) mode. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
  - *shared* (S) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to the concurrency-control manager by the transaction manager. Transaction can proceed only after request is granted.

**Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

# Atomocity Using Lock-Based Protocols

T1	T2	T3	Lock status
LOCK-S(A)			GRANT(A,T1)
READ(A)			
	LOCK-S(A)		GRANT(A,T2)
	READ(A)		
		LOCK-X(A)	WAIT(A,T3)
		WRITE(A)	

**Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

# Atomocity Using Lock-Based Protocols

T4	T5	T6	Lock status
LOCK-X(A)			GRANT(A,T4)
WRITE(A)			
	LOCK-S(A)		WAIT(A,T5)
	READ(A)		
		LOCK-X(A)	WAIT(A,T6)
		WRITE(A)	

**Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

# Atomocity Using Lock-Based Protocols

T7	T8	T9	Lock status
LOCK-X(A)			GRANT(A,T7)
WRITE(A)			
	LOCK-X(B)		GRANT(B,T8)
	WRITE(B)		
		LOCK-X(C)	GRANT(C,T9)
		WRITE(C)	

**Lock-compatibility matrix**

	S	X
S	true	false
X	false	false

# Atomocity Using Lock-Based Protocols

## Lock-compatibility matrix

**Question 24-1** Put appropriate lock and show lock status

	S	X
S	true	false
X	false	false

T1	T2	T3	T4	LOCK STATUS
READ(P)				
	WRITE(Q)			
		READ(P)		
			WRITE(P)	

# Lock-Based Protocols

- Example of a transaction performing locking:

```
 $T_2$ : lock-S( $A$ );  
      read ( $A$ );  
      unlock( $A$ );  
      lock-S( $B$ );  
      read ( $B$ );  
      unlock( $B$ );  
      display( $A+B$ )
```

- Locking as above is not sufficient to guarantee serializability — if  $A$  and  $B$  get updated in-between the read of  $A$  and  $B$ , the displayed sum would be wrong.
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.

# Pitfalls of Lock-Based Protocols

- Consider the partial schedule

$T_3$	$T_4$
lock-x ( $B$ )	
read ( $B$ )	
$B := B - 50$	
write ( $B$ )	
	lock-s ( $A$ )
	read ( $A$ )
	lock-s ( $B$ )
lock-x ( $A$ )	

- Neither  $T_3$  nor  $T_4$  can make progress — executing **lock-S( $B$ )** causes  $T_4$  to wait for  $T_3$  to release its lock on  $B$ , while executing **lock-X( $A$ )** causes  $T_3$  to wait for  $T_4$  to release its lock on  $A$ .
- Such a situation is called a **deadlock**.
  - To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back and its locks released.



# The Two-Phase Locking Protocol

- This is a protocol which **ensures conflict-serializable** schedules.
- Phase 1: Growing Phase
  - transaction may obtain locks
  - transaction may not release locks
- Phase 2: Shrinking Phase
  - transaction may release locks
  - transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e. the point where a transaction acquired its final lock).
- **Rigorous two-phase locking** is even stricter: here *all* locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.

# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions
  - Must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit** instructions as the last statement
  - By default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement

# Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency. (T1, T2, T3, ....., Tn)
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. **Conflict serializability**
  2. **View serializability**

# SERIALIZABILITY

Schedule 1 (Serial Schedule)

T1	T2
READ(A)	
READ(B)	
R=A+B	
DISPLAY(R)	
	READ(A)
	READ(C )
	A = A-1000
	C=C+1000
	WRITE(A)
	WRITE(C)

Schedule 2 (Equivalent concurrent Schedule)

T1	T2
READ(A)	
	READ(A)
	READ(C )
	A = A-1000
	C=C+1000
READ(B)	
R=A+B	
DISPLAY(R)	
	WRITE(A)
	WRITE(C)

# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- An example of a **serial** schedule in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

$A=100$ ,  $B=100$ , .In schedules 1, the sum “ $A + B$ ” is preserved.

# Schedule 2

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is **equivalent** to Schedule 1.

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit


## Question 24-2:

- Given  $A = 100$  and  $B = 100$ . Prove that the above concurrent schedule preserve database consistency.
- Explain conflict serializability

# **Concurrency Control in Distributed Databases**

# Concurrency Control

- Concurrency control mechanism maintains the ACID property of transaction by lock based protocol.
- Modify concurrency control schemes for use in distributed environment.
- We assume that each site participates in the execution of a commit protocol to ensure global transaction atomicity.
- We assume all replicas of any item are updated





# Automatic Acquisition of Locks

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:
  - if**  $T_i$  has a lock on  $D$
  - then**
  - read( $D$ )
  - else begin**
  - if necessary wait until no other
  - transaction has a **lock-X** on  $D$
  - grant  $T_i$  a **lock-S** on  $D$ ;
  - read( $D$ )
  - end**

# Automatic Acquisition of Locks (Cont.)

- **write( $D$ )** is processed as:
  - if**  $T_i$  has a **lock-X** on  $D$ 
    - then**
      - write( $D$ )
    - else begin**
      - if necessary wait until no other transaction has any lock on  $D$ ,
      - if  $T_i$  has a **lock-S** on  $D$ 
        - then**
          - upgrade** lock on  $D$  to **lock-X**
        - else**
          - grant  $T_i$  a **lock-X** on  $D$
      - write( $D$ )
    - end;**
  - All locks are released after commit or abort

# Single-Lock-Manager Approach

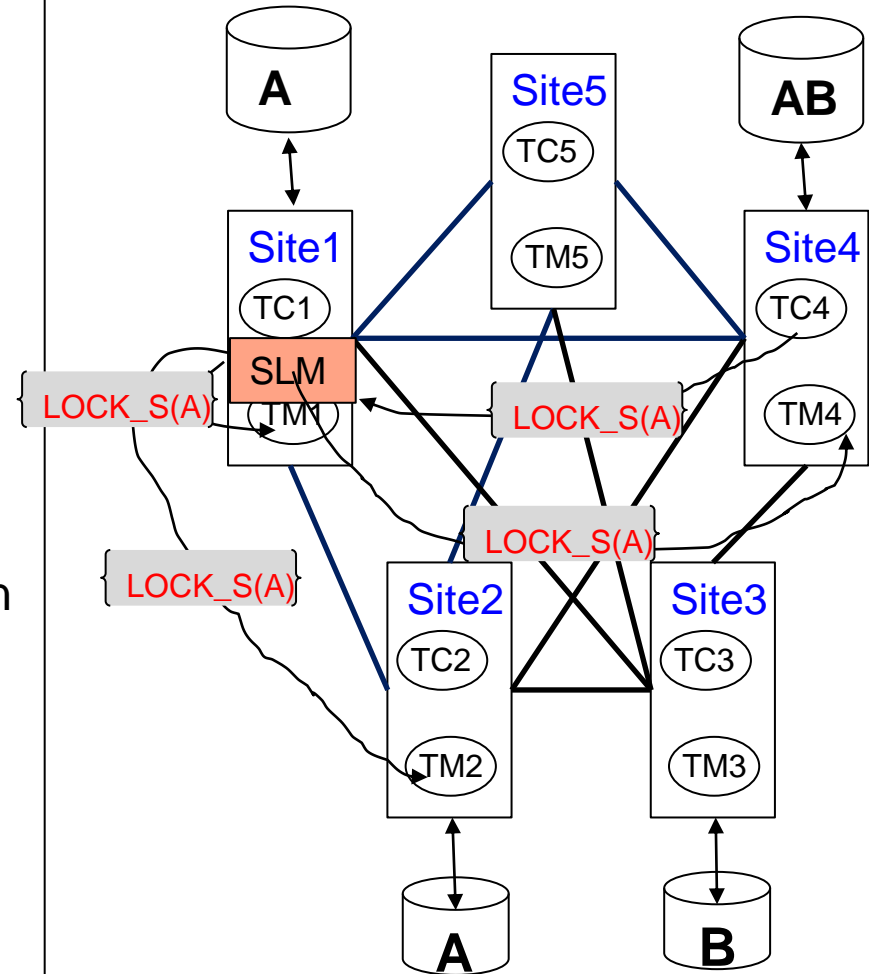
In the **single lock-manager** approach, lock manager runs on a *single* chosen site, say Site1. All lock requests sent to central lock manager

## How to perform READ (A)?

The transaction can read the data item from *any* one of the sites at which a replica of the data item resides.

**Question 25-1** Let us consider the transaction T41 at site 4. The transaction display of balance of account A. There is single lock manager at site 1.

- Write the transaction with lock.
- Write the steps to obtain the lock by T41.



# Single-Lock-Manager Approach

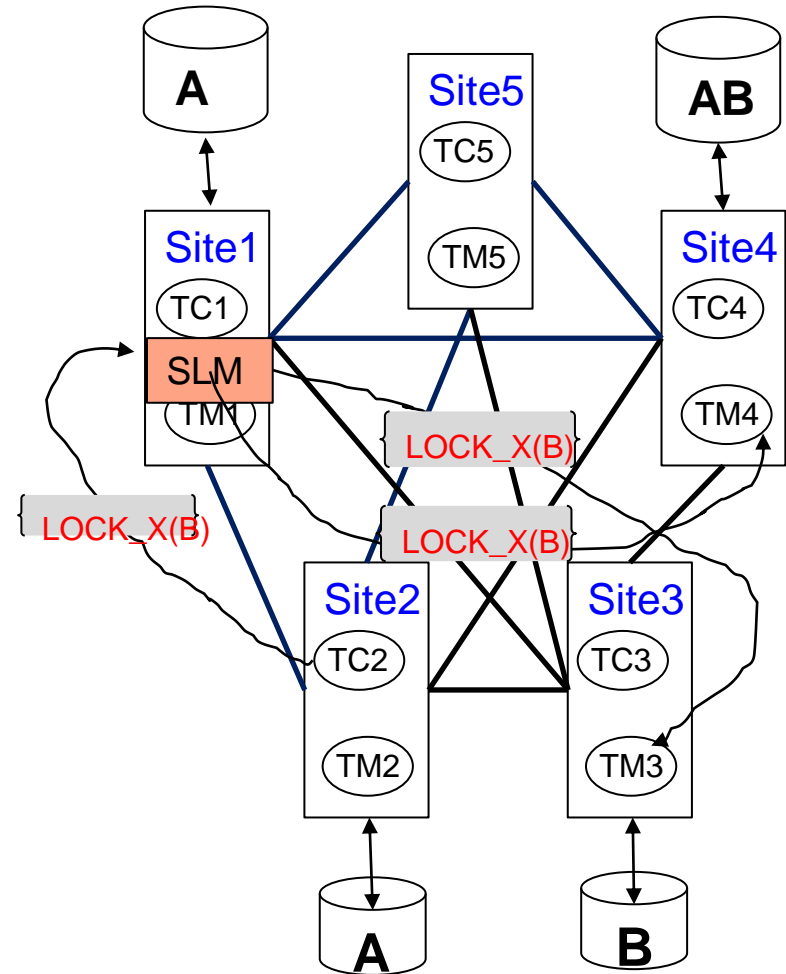
In the **single lock-manager** approach, lock manager runs on a *single* chosen site, say Site1. All lock requests sent to central lock manager

## How to perform write (B)?

The transaction has to write the data item to all of the sites at which a replica of the data item resides.

**Question 25-2** Let us consider the transaction T21 at site 2. The transaction add Tk. 1000 to account B. There is single lock manager at site 1.

- Write the transaction with lock.
- Write the steps to obtain the lock by T21.




**Writes must be performed on all replicas of a data item**

# Single-Lock-Manager Approach

- Advantages of scheme:
  - Simple implementation
  - Simple deadlock handling
- Disadvantages of scheme are:
  - Bottleneck: lock manager site becomes a bottleneck
  - Vulnerability: system is vulnerable to lock manager site failure.

# Distributed Lock Manager

- In the **distributed lock-manager** approach, functionality of locking is implemented by lock managers at each site
- Lock managers control access to local data items
- Locking is performed separately on each site accessed by transaction
- **Every replica** must be locked and updated
- But special protocols may be used for replicas (more on this later)



# Distributed Lock Manager


Advantage: work is distributed and can be made robust to failures

Disadvantage:

Possibility of a global deadlock without local deadlock at any single site

Lock managers must cooperate for deadlock detection

Compare single lock manager and distributed lock manager approach.



# Replication

- **High availability** is a key goal in a distributed database
  - **Robustness**: the ability to continue function despite failures
- Replication is key to robustness
- Replication decisions can be made at level of data items, or at the level of partitions



# Concurrency Control With Replicas

Focus here on concurrency control with locking

Failures addressed later

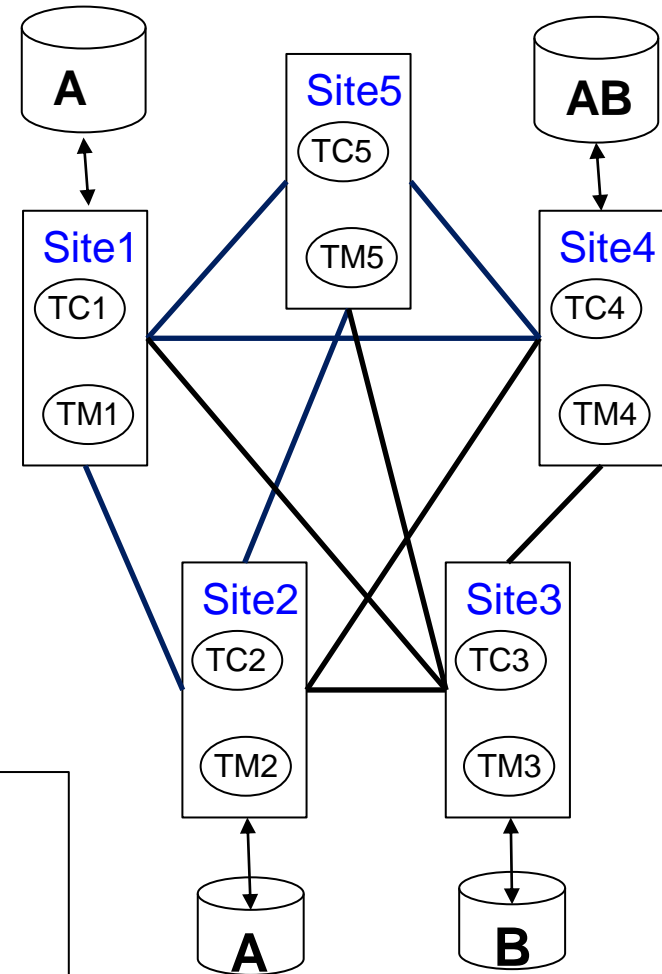
Ideas described here can be extended to other protocols

## Primary copy

- one replica is chosen as primary copy for each data item (e.g., A in site 2)
- Node containing primary replica is called **primary node (site 2 for A)**
- concurrency control decisions made at the primary copy only (by site 2)

**Example:** How LOCK-X(A) of T10 at site 3 will be obtained using primary protocol? Site 2 is primary copy for A.

1. site 3 sends LOCK-X(A) request to sites 2.
2. If T10 gets lock on A, then A is locked and T10 proceeds



# Concurrency Control With Replicas

Focus here on concurrency control with locking

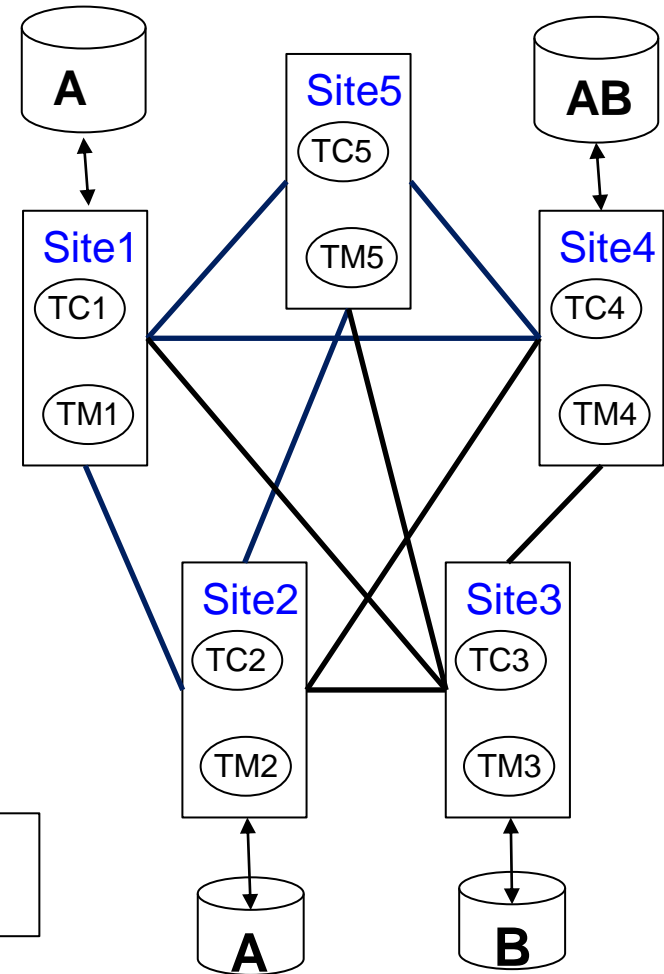
Failures addressed later

Ideas described here can be extended to other protocols

## Primary copy


- one replica is chosen as primary copy for each data item (e.g., A in site 2)
- Node containing primary replica is called **primary node (site 2 for A)**
- concurrency control decisions made at the primary copy only (by site 2)

**Question:** How LOCK-X(B) of T2 at site 2 will be obtained using primary copy protocol? Site 4 is primary copy for B.



# Concurrency Control With Replicas

- Benefit: Low overhead of locking (How?)
- Drawback: primary copy failure (site 2 for A) results in loss of lock information and non-availability of data item (A), even if other replicas (site 1, site 4) are available
- Extensions to allow backup server to take over possible, but vulnerable to problems on network partition



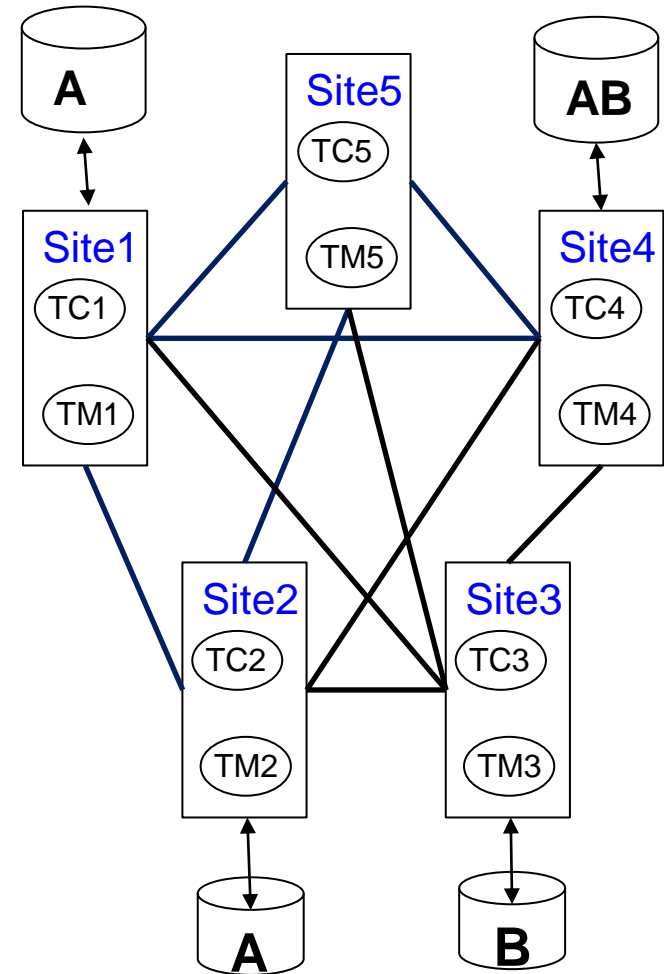
# Concurrency Control With Replicas (Cont.)

## Majority protocol:

- Transaction requests locks at majority/all replicas
- Lock is successfully acquired on the data item only if lock obtained at a majority of replicas

**Example:** How LOCK-X(A)/ LOCK-S(A) of T10 at site 3 will be obtained using majority protocol?

1. site 3 sends LOCK-X(A)/ LOCK-S(A) request to sites 1, 2 or 2,4 or 1,4 or 1,2,4.
2. If T1 gets lock on 2 sites, then A is locked and T1 proceeds



# Concurrency Control With Replicas (Cont.)

## Number of messages in Majority Protocol

### Benefit:

- Resilient to node failures
- Processing can continue as long as at least a majority of replicas are accessible

### Overheads

- Higher cost due to multiple messages
- Possibility of deadlock even when locking single item

**Example:** A data item Q is replicated to 10 sites. Transaction T1 has LOCK-S(A).

What is the minimum number of messages required to obtain this lock using majority protocol?

Messages for lock request  
 $= 10/2 + 1 = 6$

Messages for lock grant  
 $= 10/2 + 1 = 6$

Messages for unlock  
 $= 10/2 + 1 = 6$

Total =  $6+6+6 = 18$

# Concurrency Control With Replicas (Cont.)

## Number of messages in Majority Protocol

### Benefit:

- Resilient to node failures and node failures
- Processing can continue as long as at least a majority of replicas are accessible

### Overheads

- Higher cost due to multiple messages
- Possibility of deadlock even when locking single item

**Question 26-1:** A data item P is replicated to 12 sites. Transaction T2 has LOCK-X(P).

Find the minimum number of messages required to

- a. obtain this lock using majority protocol?
- b. unlock

# Concurrency Control With Replicas (Cont.)

## Number of messages in Majority Protocol

### Benefit:

- Resilient to node failures and node failures
- Processing can continue as long as at least a majority of replicas are accessible

### Overheads

- Higher cost due to multiple messages
- Possibility of deadlock even when locking single item

How can you avoid such deadlocks?

A data item Q is replicated to  $n$  sites. Transaction T1 has LOCK-S(A).

What is the minimum number of messages required to obtain this lock?

Messages for lock request  
 $= n/2 + 1$

Messages for lock grant  
 $= n/2 + 1$

Messages for unlock  
 $= n/2 + 1$

Total  $= 3(n/2 + 1)$

# Concurrency Control With Replicas (Cont.)

## Number of messages in Primary Copy Protocol

**Example:** A data item Q is replicated to 10 sites. Transaction T1 has LOCK-S(A). What is the number of messages required to obtain this lock using primary copy protocol?

Messages for lock request  
= 1

Messages for lock grant  
= 1

Messages for unlock  
= 1

Total =  $1+1+1 = 3$



# Concurrency Control With Replicas (Cont.)

## Biased protocol

- Shared lock can be obtained on any replica
- Reduces overhead on reads
- Exclusive lock must be obtained on *all* replicas
- Blocking if any replica is unavailable

**Example:** A data item Q is replicated to 10 sites. Transaction T1 has **LOCK-X(A)**.

What is the minimum number of messages required to obtain this lock?

No. of messages =  $3 \times 10 = 30$

**Example:** A data item Q is replicated to 10 sites. Transaction T1 has **LOCK-S(A)**.

What is the minimum number of messages required to obtain this lock?

Messages for lock request

= 1

Messages for lock grant

= 1

Messages for unlock

= 1

Total =  $1 + 1 + 1 = 3$

# Concurrency Control With Replicas (Cont.)

## Biased protocol

- Shared lock can be obtained on any replica
- Reduces overhead on reads
- Exclusive lock must be obtained on *all* replicas
- Blocking if any replica is unavailable

**Question26-2:** Data item A and B are replicated to 15 sites. Transaction T11 has LOCK-X(A) and LOCK-S(B).

- a. What are the minimum number of sites required to obtain these locks using biased protocol?
- b. Find the number of messages to obtain lock and release lock in each case.
- c. Compare biased protocol with majority protocol in terms of performance.

# Quorum Consensus Protocol

**Quorum consensus** protocol for locking  
Each site is assigned a weight; let  $S$  be the total of all site weights

Choose two values **read quorum**  $Q_R$  and **write quorum**  $Q_W$

Such that  $Q_R + Q_W > S$  and  $2 * Q_W > S$

Each read must lock enough replicas that the sum of the site weights

is  $\geq Q_R$

Each write must lock enough replicas that the sum of the site weights

is  $\geq Q_W$

Can choose  $Q_R$  and  $Q_W$  to tune relative overheads on reads and writes

Suitable choices result in majority and biased protocols.

**What are they?**

**Example:** The data item  $P$  is replicated in 9 sites and the weight of each site is 1.  $Q_W$  is 7 and  $Q_R$  is 3.  $S=9$

Find the number of replicas to be locked for

- i. LOCK-S
- ii. LOCK-X

**For Lock-S,  $Q_r = 3$ ,**

Sum of site weight = 3

Number of sites = 3

**For Lock-X,  $Q_w = 7$**

Sum of site weight = 7

Number of sites = 7

# Quorum Consensus Protocol

**Quorum consensus** protocol for locking  
Each site is assigned a weight; let  $S$  be the total of all site weights  
Choose two values **read quorum**  $Q_R$  and **write quorum**  $Q_W$   
Such that  $Q_R + Q_W > S$  and  $2 * Q_W > S$   
Each read must lock enough replicas that the sum of the site weights is  $\geq Q_R$   
Each write must lock enough replicas that the sum of the site weights is  $\geq Q_W$   
Can choose  $Q_R$  and  $Q_W$  to tune relative overheads on reads and writes  
Suitable choices result in majority and biased protocols.  
**What are they?**

**Example:** The data item  $P$  is replicated in 9 sites  $S_1, S_2, S_3, S_4, S_5, S_6, S_7, S_8, S_9$  and the weight of the sites are 1,1,1,2,2,2,3,3,3 respectively.  $Q_W$  is 12 and  $Q_R$  is 7.

Find the replicas to be locked for

- i. LOCK-S( $P$ )
- ii. LOCK-X( $P$ )
  - a. Starting the sites with lowest weight with ascending order.
  - b. Starting the sites with highest weight with descending order.

## a. Ascending order

Weights  $S_1 + S_2 + S_3 + S_4 + S_5$   
 $= 1+1+1+2+2 = 7 = Q_R$

For LOCK-S( $P$ ), sites to have locked is  $S_1, S_2, S_3, S_4, S_5$

Weights  $S_1+S_2+S_3+S_4+S_5+S_6+S_7$   
 $= 1+1+1+2+2+2+3 = 12 = Q_W$

For LOCK-X( $P$ ), sites to have locked is  $S_1, S_2, S_3, S_4, S_5, S_6, S_7$

# Quorum Consensus Protocol

**Quorum consensus** protocol for locking  
Each site is assigned a weight; let  $S$  be the total of all site weights  
Choose two values **read quorum**  $Q_R$  and **write quorum**  $Q_W$   
Such that  $Q_R + Q_W > S$  and  $2 * Q_W > S$   
Each read must lock enough replicas that the sum of the site weights is  $\geq Q_R$   
Each write must lock enough replicas that the sum of the site weights is  $\geq Q_W$   
Can choose  $Q_R$  and  $Q_W$  to tune relative overheads on reads and writes  
Suitable choices result in majority and biased protocols.

**What are they?**

**Example:** The data item  $P$  is replicated in 9 sites  $S1, S2, S3, S4, S5, S6, S7, S8, S9$  and the weight of the sites are 1,1,1,2,2,2,3,3,3 respectively.  $Q_W$  is 12 and  $Q_R$  is 7.

Find the replicas to be locked for

- i. LOCK-S( $P$ )
- ii. LOCK-X( $P$ )
  - a. Starting the sites with lowest weight with ascending order.
  - b. Starting the sites with highest weight with descending order.

## **b. Descending order**

Weights  $S9+S8+S7$   
 $= 3+3+3 = 9 > Q_R$

For LOCK-S( $P$ ), sites to have locked are  $S9, S8, S7$

Weights  $S9+S8+S7+S6+S5$   
 $= 3+3+3+2+2 = 13 > Q_W$

For LOCK-X( $P$ ), sites to have locked are  $S9, S8, S7, S6, S5$

# Quorum Consensus Protocol

**Quorum consensus** protocol for locking  
Each site is assigned a weight; let  $S$  be the total of all site weights

Choose two values **read quorum**  $Q_R$  and **write quorum**  $Q_W$

Such that  $Q_R + Q_W > S$  and  $2 * Q_W > S$

Each read must lock enough replicas that the sum of the site weights is  $\geq Q_R$

Each write must lock enough replicas that the sum of the site weights is  $\geq Q_W$

Can choose  $Q_R$  and  $Q_W$  to tune relative overheads on reads and writes

Suitable choices result in majority and biased protocols.

**What are they?**

**Question:** The data item  $P$  is replicated in 8 sites  $S_1, S_2, S_3, S_4, S_5, S_6, S_7, S_8$  and the weight of the sites are 2,2,2,3,3,3,4,4 respectively.

a. Find  $Q_W$  and  $Q_R$ .

Find the replicas to be locked for

- LOCK-S( $P$ )
- LOCK-X( $P$ )

b. Starting the sites with lowest weight with ascending order.

c. Starting the sites with highest weight with descending order.